



gkwdist :: cheat sheet

Seven **nested** distributions for bounded continuous data on the open interval (0, 1) — proportions, rates, shares, indices. Densities, CDFs, quantiles, random generation, and **analytical** log-likelihood, score and Hessian, all in C++.


version 1.1.6

evandeilton.github.io/gkwdist

install.packages("gkwdist")

THE FAMILY one distribution, seven models

Every model below is **GKw with parameters held fixed**. Follow an edge to fix a parameter and drop to a simpler model; the number in the badge is how many free parameters remain.



Closed-form CDF
and quantile:
Kw · EKw · KKw

Uniform is GKw(1, 1, 1, 0, 1) — the neutral value of every parameter. **Kw** and **Beta** are *not* nested in one another, and neither are **EKw**/**Mc** or **BKw**/**KKw**.

WHAT EACH PARAMETER DOES

$\alpha > 0$	Exponent on x. Governs the density near 0.
$\beta > 0$	Exponent on $1-x$. Governs the density near 1.
$\gamma > 0$	The Beta(γ , $\delta+1$) generator that reshapes the Kumaraswamy core. $\delta = 0$ — not 1 — is the neutral value.
$\delta \geq 0$	
$\lambda > 0$	Exponentiation (power) parameter.

Boundary rule — what happens at x = 0 and x = 1

Each endpoint is decided by the sign of a single exponent. `d*( )` returns the **limit**, as base R does, not 0.

AT	EXPONENT	> 0	= 0	< 0
x=0	$\alpha \cdot \gamma \cdot \lambda - 1$	0	const	Inf
x=1	$\beta \cdot (\delta + 1) - 1$	0	const	Inf

Each sub-family inherits the rule through its own fixed values: for Kw(α, β) they read $\alpha-1$ and $\beta-1$. `dkw(0, 0.5, 1)` is Inf; `dbeta(1, 2, 0)` is 2.

49 FUNCTIONS, ONE NAMING RULE

prefix + family code. Row = your model, column = the job. **blue** distribution API **red** likelihood API.

FAMILY (code)	D PDF	P CDF	Q QUANT	R RAND	LL -LOGL	GR -SCORE	HS -HESS
Gen. Kumaraswamy gkw · 5p	dgkw	pgkw	qgkw	rgkw	llgkw	grgkw	hsgkw
Beta-Kumaraswamy bkw · 4p	dbkw	pbkw	qbkw	rbkw	llbkw	grbkw	hsbkw
Kum.-Kumaraswamy kkw · 4p	dkkw	pkkw	qkkw	rkkw	llkkw	grkkw	hskkw
Exp. Kumaraswamy ekw · 3p	dekwx	pekwx	qekwx	rekwx	llekw	grekw	hsekw
McDonald mc · 3p	dmc	pmc	qmc	rmc	llmc	grmc	hsmc
Kumaraswamy kw · 2p	dkw	pkw	qkw	rkwx	llkw	grkw	hskw
Beta beta_ · 2p	dbeta_	pbeta_	qbeta_	rbeta_	llbeta_	grbeta_	hsbeta_

Mind the underscore. Only the Beta row is irregular: `d/p/q/r` keep the trailing `_` so they do not mask stats: `dbeta` — the likelihood trio does not. `llbeta`, not `llbeta_`.

DISTRIBUTION FUNCTIONS `d · p · q · r`

```
dgkw(x, alpha, beta, gamma, delta, lambda, log = FALSE)
pgkw(q, ..., lower.tail = TRUE, log.p = FALSE)
qgkw(p, ..., lower.tail = TRUE, log.p = FALSE)
rgkw(n, ...)
```

A sub-family drops the parameters it fixes: `dkw(x, alpha, beta, log = FALSE)`.


log, log.p	Work on the log scale. Use them instead of <code>log(d...)</code> near the boundaries.
lower.tail	FALSE gives $P(X > q)$, computed directly rather than as $1 - p$.
n	<code>length(n) > 1</code> means <code>length(n)</code> draws; <code>n = 0</code> gives <code>numeric(0)</code> .

The base-R contract these keep

- Recycling.** x and every parameter are recycled to a common length: `dkw(c(.3,.6), c(1,2,3,4), 2)` gives 4 values.
- Attributes.** `dim`, `dimnames` and names of the first argument survive when the lengths agree, as in `stats::dbeta`.
- Outside [0,1]** the density is 0; at the two endpoints it is the limit (box, left).
- An invalid parameter is an error** — $\alpha \leq 0$, $\delta < 0$, NA all stop(). They do not return NaN.
- Defaults are the neutral values**, so `dgkw(x)` is the uniform density.

SHAPES THE FAMILY REACHES

Real output. The boundary rule (column 1) predicts each end.




Unimodal
dkw(x, 2, 5)



U-shaped
dkw(x, 0.5, 0.5)



J — increasing
dkw(x, 3, 0.6)



Reverse J
dkw(x, 0.6, 3)



Spike at 0 + mode
dgkw(x, .2, .4, .6, 5, 6)



Mode + spike at 1
dgkw(x, 6, .2, 1.2, 2, .2)

What it will not do. At most **one interior mode** — a second peak is always a divergence at 0 or 1, as in the bottom row. Two interior humps need a mixture, not a bigger GKw.

Where this Beta differs from stats

$\delta = 0$ is neutral, so the second shape sits at $\delta + 1$:

```
dbeta_(x, gamma, delta) == stats::dbeta(x, gamma, delta + 1)
```

That shape is therefore always ≥ 1 : this Beta **cannot be U-shaped**. Use Kw, EKw or GKw for a bathtub on (0,1).

PICK A STARTING FAMILY

YOUR DATA	START WITH
Single hump, roughly symmetric	Beta 2p
Single hump, skewed	Kw 2p — closed-form CDF & quantile
Monotone (J or reverse J)	Kw or Beta
U-shaped / mass at both ends	Kw, EKw ($\alpha, \beta < 1$)
Heavy tail on one side only	EKw or Mc 3p
An explicit power transform	Mc λ is that power
Stubborn residual shape	BKw, KKw, then GKw

Fit up, not down. Add a parameter only when a likelihood-ratio test (page 2) says it earns its place.

FIT BY MAXIMUM LIKELIHOOD

Objective, gradient and Hessian drop straight into `stats::optim`.

```
x <- rekw(2000, 2, 3, 1.5)      # data on (0,1)

start <- gkwgetstartvalues(x, family = "ekw")

fit <- optim(start,
  fn      = llekwx,  # -logL
  gr      = grekwx,  # -score
  data    = x,      # passed through
  method  = "BFGS",
  hessian = TRUE)

est <- fit$par                # MLE
logLik <- -fit$value          # note the sign
se  <- sqrt(diag(solve(fit$hessian)))
ci  <- est + outer(se, c(-1.96, 1.96))
```

`fn` is already the *negative* log-likelihood, so `optim`'s default minimisation is the right direction and `fit$hessian` is the **observed information** — invert it directly, no sign flip.

For the exact information matrix, use the analytical Hessian:

```
se <- sqrt(diag(solve( hsekw(est, x) )))
```

They agree to ~5e-7 relative: `optim`'s is a finite difference.

STARTING VALUES

```
gkwgetstartvalues(x, family = "gkw", n_starts = 5)
```

Method of moments by Nelder–Mead, returning a **named** vector already in the par order that family's `ll*` expects. family is one of "gkw" "bkwx" "kkwx" "ekwx" "mc" "kw" "beta" (case-insensitive).

- **Deterministic.** `set.seed()` has no effect and it never touches `.Random.seed`. Widen the search with `n_starts`; there is no seed.
- Four fixed starting points are always used, so `n_starts < 4` behaves as 4; cost grows linearly and the objective is non-increasing in `n_starts`.
- Results are clipped to a box — $\alpha, \beta \in (0.1, 50)$, $\gamma \in (0.1, 10)$ but $(0.1, 50)$ for "beta", $\delta \in (0.01, 10)$, $\lambda \in (0.1, 20)$ — so a value sitting exactly on a bound is a **hint that the box bit**, not an estimate.
- Values outside $(0,1)$ are truncated **with a warning**. Treat that warning as a scale problem in the data — a 0–100 percentage — not as noise.

RETURN VALUES UNDER FAILURE

SITUATION	LL*	GR* / HS*
Bad par (≤ 0 , wrong length, NA)	Inf	NaN
Data outside (0,1)	Inf + warning	NaN

An infinite objective gives the optimiser no direction — it sits there reporting convergence. **Screen first:** `stopifnot(all(x > 0 & x < 1))`. An exact 0 or 1 is the usual culprit; squeeze it in with $(x*(n-1) + 0.5)/n$.

LIKELIHOOD FUNCTIONS $ll \cdot gr \cdot hs$

```
llgkw(par, data)  ~ scalar  -ℓ(θ)
grgkw(par, data)  ~ vector  -∇ℓ(θ)
hsgkw(par, data)  ~ matrix  -∇²ℓ(θ)
```

All three carry a **minus sign**: minimise `ll*`, and `hs*(est, x)` is the observed information. `par` is a plain numeric vector — the names are ignored, only the **order** matters.

⚠ **par order differs by family**

The vector is **positional**. A mis-ordered `par` is not an error — it silently fits the wrong model.

FAMILY	LEN	PAR = C(...)
gkw	5	alpha, beta, gamma, delta, lambda
bkwx	4	alpha, beta, gamma, delta — no λ
kkwx	4	alpha, beta, delta, lambda — no γ
ekwx	3	alpha, beta, lambda
mc	3	gamma, delta, lambda
kw	2	alpha, beta
beta	2	gamma, delta

bkwx and **kkwx** are the trap: both take four, and they are not the same four. Let `gkwgetstartvalues()` build the vector and the order comes out right for free.

WHY ANALYTICAL DERIVATIVES

One pass over the data instead of a finite-difference stencil per parameter. Measured at `n = 20 000`, `GKw`, against `numDeriv`:

	GR*(θ)	NUMDERIV	
score	2.8 ms	61 ms	≈22×
Hessian	4.0 ms	177 ms	≈44×

Agreement with `numDeriv` to ~9e-6 absolute — the difference is the finite-difference truncation error, not the analytical form.

FIT ALL SEVEN, THEN RANK

```
fam <- c("gkw", "bkwx", "kkwx", "ekwx", "mc", "kw", "beta")
ll  <- list(llgkw, llbkwx, llkkwx, llekwx, llmc, llkw, llbeta)
gr  <- list(grgkw, grbkwx, grkkwx, grekw, grmc, grkw, grbeta)

tab <- Map(function(f, l, g) {
  s <- gkwgetstartvalues(x, family = f)
  o <- optim(s, l, g, data = x, method = "BFGS")
  c(k = length(s), AIC = 2*o$value + 2*length(s))
}, fam, ll, gr)

tab <- do.call(rbind, tab); tab[order(tab[, "AIC"]), ]
```

The three lists run in the same order as page 1's matrix, and `gkwgetstartvalues()` returns each `par` already in the order its family expects — nothing is aligned by hand.

COMPARE NESTED MODELS

Comparable by likelihood-ratio test only if one is the other with parameters fixed; `df` counts them.

FULL	REDUCED	CONSTRAINT	DF
gkw	bkwx	$\lambda = 1$	1
gkw	kkwx	$\gamma = 1$	1
gkw	mc	$\alpha = \beta = 1$	2
bkwx	beta	$\alpha = \beta = 1$	2
kkwx	ekwx	$\delta = 0 \uparrow$	1
mc	beta	$\lambda = 1$	1
ekwx	kw	$\lambda = 1$	1
gkw	ekwx	$\gamma = 1, \delta = 0 \uparrow$	2
gkw	kw	$\gamma = 1, \delta = 0, \lambda = 1 \uparrow$	3
gkw	beta	$\alpha = \beta = \lambda = 1$	3

```
lrt <- 2 * (fit_red$value - fit_full$value)
pchisq(lrt, df, lower.tail = FALSE)
```

$\uparrow \delta = 0$ is on the **boundary**, so χ^2 (df) is conservative: the null is $0.5 \cdot \chi^2$ (df-1) + $0.5 \cdot \chi^2$ (df).

Not nested: kw vs beta, ekwx vs mc, bkwx vs kkwx. Compare those with AIC/BIC only.

INFORMATION CRITERIA

```
k <- length(fit$par); ll <- -fit$value
AIC <- -2*ll + 2*k
BIC <- -2*ll + k*log(length(x))
```

Lower is better; a gap under ~2 is not evidence. All seven are fitted to the *same* `x`, so they compare directly.

CHECK THE FIT

```
# PIT residuals — uniform if the model holds
u <- pekwx(x, est[1], est[2], est[3])
ks.test(u, "punif")

# Q-Q against the fitted quantiles
q <- qekwx(ppoints(length(x)), est[1], est[2], est[3])
qqplot(q, sort(x)); abline(0, 1, col = 2)
```

Estimating from the same `x` makes the KS p-value optimistic.

Traps worth a second look

- $\delta = 0$, not 1, is neutral — `dgkw(x, 1, 1, 1, 0, 1)` is uniform.
- `llbeta` has no underscore, `dbeta_` does — and `dbeta_(x, g, d)` is `stats::dbeta(x, g, d+1)`.
- `ll*` returns +Inf, never an error, on a contaminated sample.
- `gkwgetstartvalues()` ignores `set.seed()`.
- `%>%` is re-exported but **deprecated** — take it from `magrittr`, or use `|>`.